

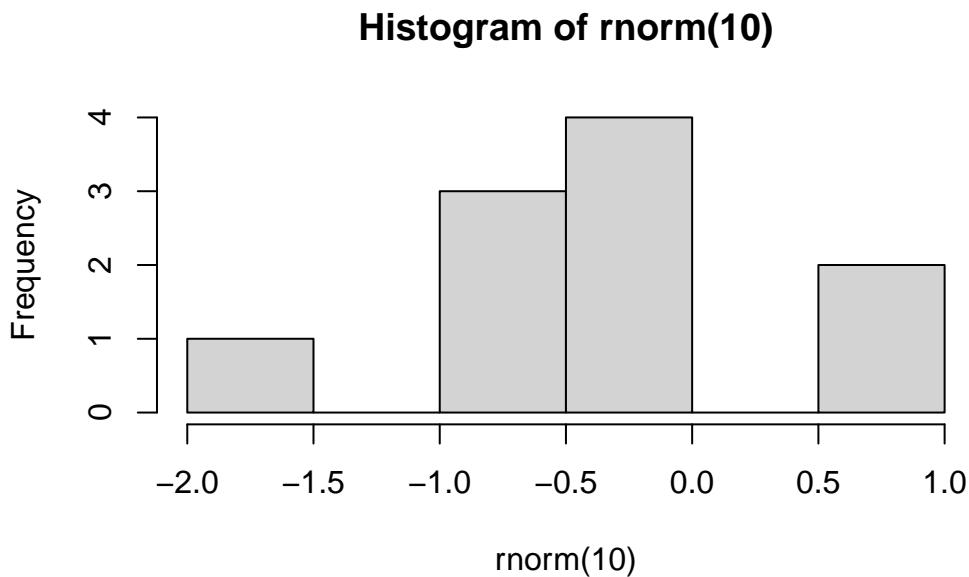
Class07:Machine Learning1

Ebru Robinson

Today we will begin our exploration of some “classical” machine learning approaches. We will start with clustering:

Let’s first make up some data to cluster where we know what the answer should be.

```
hist(rnorm(10))
```



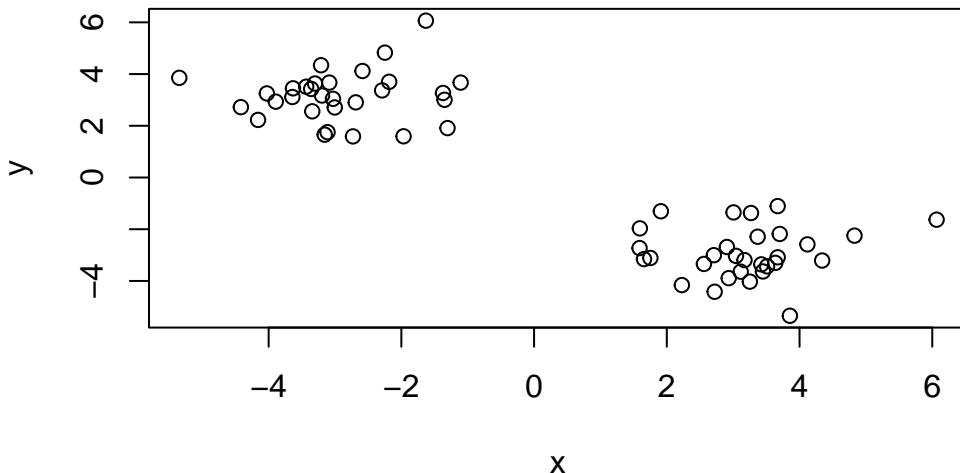
```
x <- c(rnorm(30,mean=-3),rnorm(30,mean=3))
y <- rev(x)

x <- cbind(x,y)
head(x)
```

```
x      y  
[1,] -3.030354 3.042981  
[2,] -3.112372 1.752562  
[3,] -1.966876 1.593968  
[4,] -3.155117 1.656351  
[5,] -5.346527 3.854675  
[6,] -3.342669 2.559018
```

a wee peak at

```
plot(x)
```



the main function in “base” R for k-means clustering is called `kmeans()`.

```
k <- kmeans(x, centers=4)  
k
```

K-means clustering with 4 clusters of sizes 12, 10, 20, 18

Cluster means:

```
x      y  
1  2.594342 -2.188377
```

```
2 -1.928161  3.827909
3 -3.426093  2.838432
4  3.550869 -3.419053
```

Clustering vector:

```
[1] 3 3 3 3 3 3 3 2 3 3 3 2 3 2 2 2 3 3 2 3 3 2 2 2 3 2 3 3 3 1 4 4 4 1 4 1 1
[39] 1 4 1 4 4 4 4 1 4 1 4 4 4 1 4 4 4 4 1 1 1 4
```

Within cluster sum of squares by cluster:

```
[1] 14.15911 15.44459 19.85046 25.12928
(between_SS / total_SS =  93.9 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q. How big are the clusters(i.e their size)?

```
k$size
```

```
[1] 12 10 20 18
```

Q. What clusters do my data points reside in?

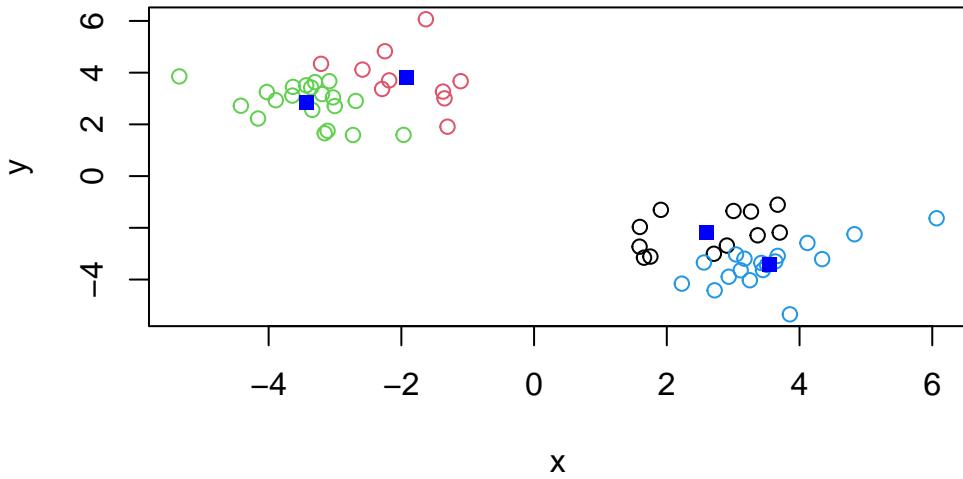
```
k$cluster
```

```
[1] 3 3 3 3 3 3 3 2 3 3 3 2 3 2 2 2 3 3 2 3 3 2 2 2 3 2 3 3 3 1 4 4 4 1 4 1 1
[39] 1 4 1 4 4 4 4 1 4 1 4 4 4 1 4 4 4 4 1 1 1 4
```

Q. Make a plot of our data colored by cluster assignment- i.e. make a result figure...

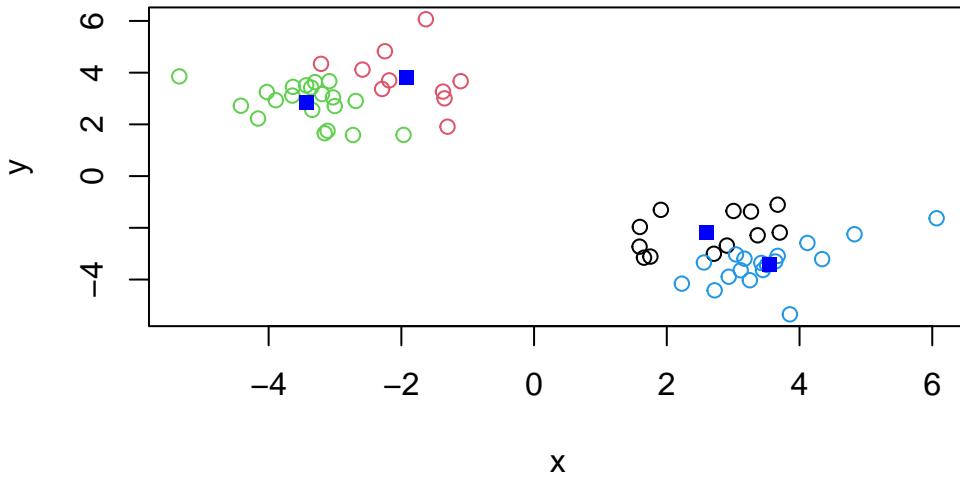
```
plot(x, col=k$cluster)
points(k$centers, centers=4, col="blue", pch=15)
```

Warning in plot.xy(xy.coords(x, y), type = type, ...): "centers" is not a graphical parameter



```
k4 <- kmeans(x,centers=4)
plot(x, col=k$cluster)
points(k$centers,centers=4,col="blue", pch=15)
```

Warning in plot.xy(xy.coords(x, y), type = type, ...): "centers" is not a graphical parameter



Q. Run kmeans with values center(i.e values of k) equal 1 to 6

```
k1 <- kmeans(x,centers=1)$tot.withinss
k2 <- kmeans(x,centers=2)$tot.withinss
k3 <- kmeans(x,centers=3)$tot.withinss
k4 <- kmeans(x,centers=4)$tot.withinss
k5 <- kmeans(x,centers=5)$tot.withinss
k6 <- kmeans(x,centers=6)$tot.withinss
ans <- c(k1,k2,k3,k4,k5,k6)
```

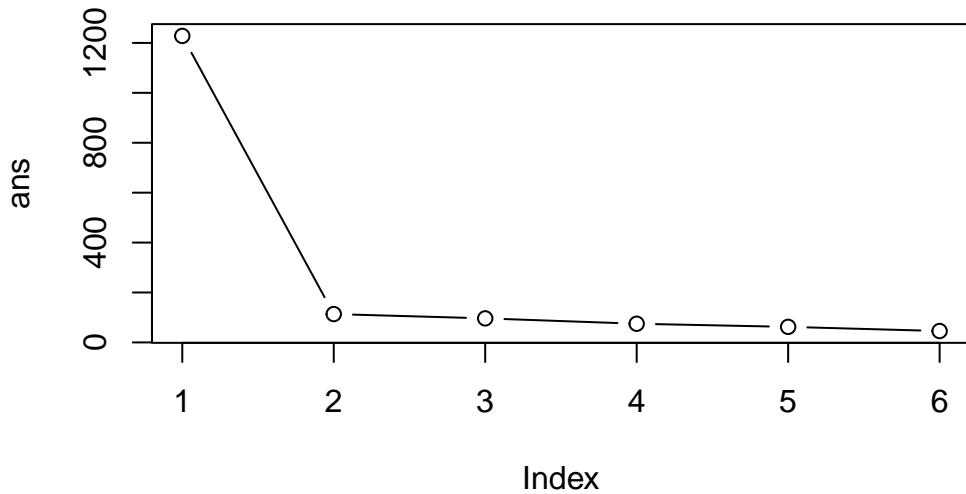
or use a for loop

```
ans <- NULL
for(i in 1:6) {
  ans <- c(ans, kmeans(x,centers=i)$tot.withinss)
}
ans
```

```
[1] 1228.04720 113.56166 96.32075 74.83497 62.34236 45.60488
```

Make a “scree-plot”

```
plot(ans, typ="b")
```



##Hierarchical Clustering

The main function in “base” R for this is called `hclust()`

```
d <- dist(x)
hc <- hclust(d)
hc
```

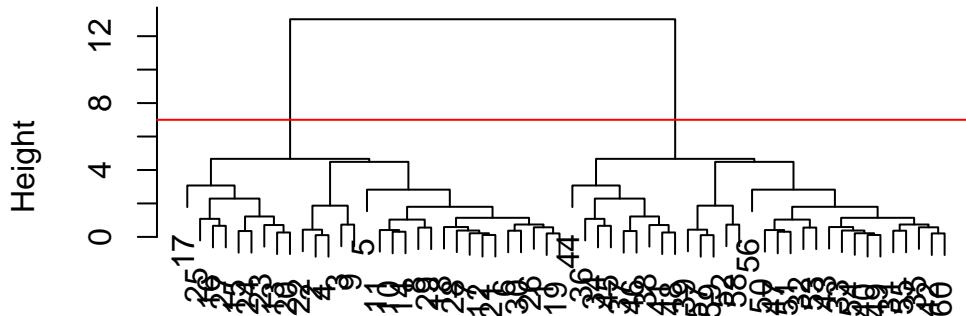
Call:

`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc)
abline(h=7, col="red")
```

Cluster Dendrogram



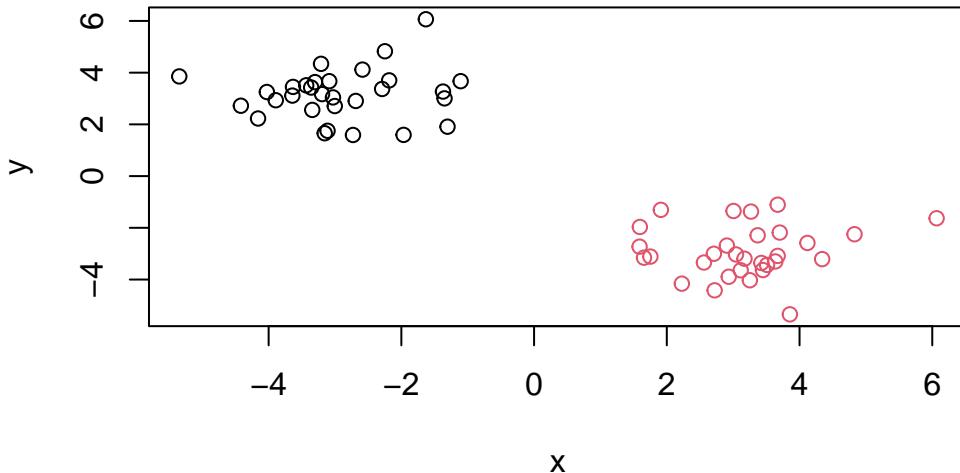
```
d  
hclust (*, "complete")
```

to obtain clusters from our `hclust` result object `hc` we “cut” the tree to yield different sub branches. For this we use the `cutree()` function

```
grps <- cutree(hc, h=7)  
grps
```

Results figure

```
plot(x,col=grps)
```



```
#install.packages("pheatmap")
```

```
library("pheatmap")
```

Principal Component Analysis (PCA)

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this question?

This will return both the number of rows and columns in the data frame x.

```
dim(x)
```

```
[1] 17 5
```

```
## Preview the first 6 rows  
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
# Note how the minus indexing works  
rownames(x) <- x[,1]  
x <- x[,-1]  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
x <- read.csv(url, row.names = 1)  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

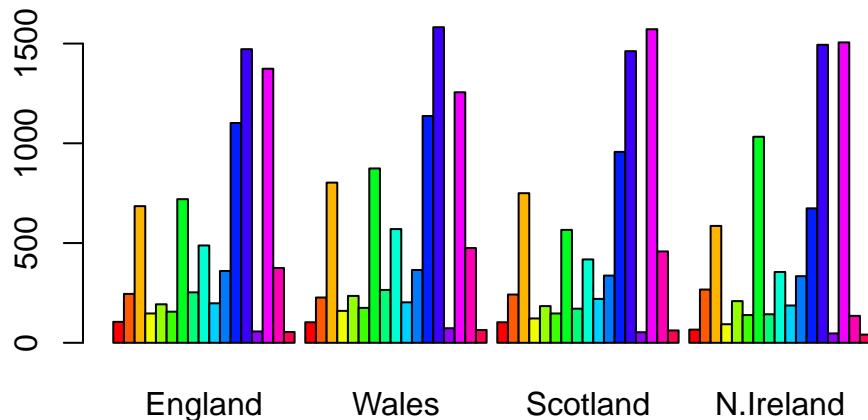
Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I like fixing it up front when reading the data...

I prefer the import-time approach: set row names when you read the file, e.g. `read.csv(url, row.names = 1)`. It is more robust and reproducible.

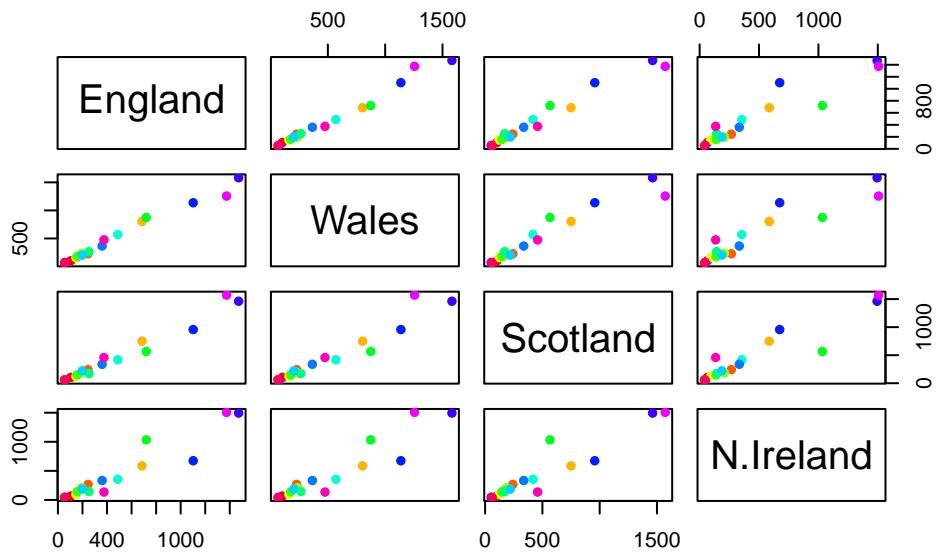
`##spotting major differences and trends`

```
# Using base R  
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

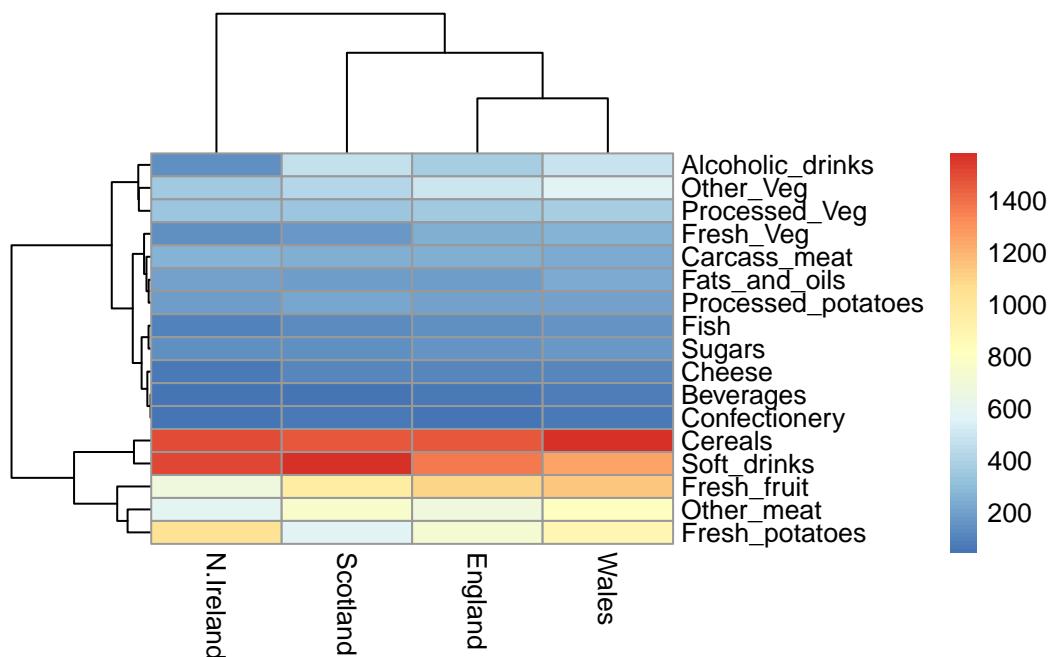


`## Pairs plots and heatmaps` Scatterplot matrices (“pairs plots”) can be useful for relatively small datasets like this one. lets have a look:

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

It looks like Wales and England are quite similar in their consumption of these foods. It is still quite difficult to tell what is going on in the dataset

#PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`.

As we want to do PCA on the food data for the different countries we will want the foods in the columns

```
# Use the prcomp() PCA function
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Our result object is called `pca` and it has a `$x` component that will look at first

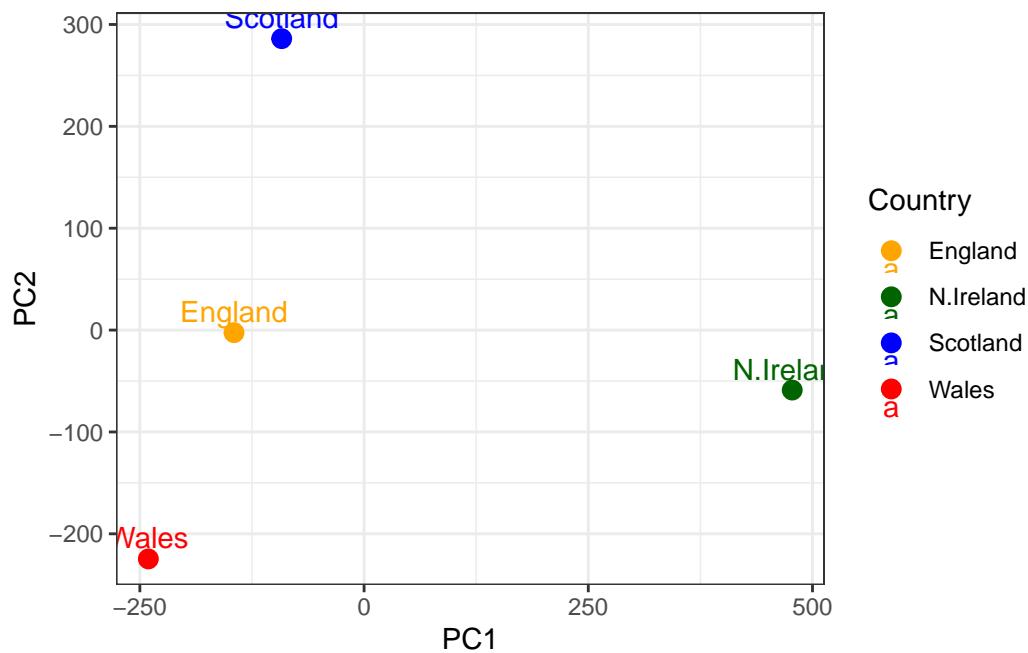
```
library(ggplot2)

# Convert PCA results to a data frame
df <- as.data.frame(pca$x)
df$Country <- rownames(df)

# Define colors in the same order as your countries appear in df
cols <- c("orange", "red", "blue", "darkgreen")
names(cols) <- df$Country # optional: name colors explicitly for safety

# Plot
ggplot(df, aes(x = PC1, y = PC2, label = Country, color = Country)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  scale_color_manual(values = cols) +
  xlab("PC1") +
```

```
ylab("PC2") +  
theme_bw()
```



Another major result out of PCA is the so-called “variable” loadings or \$rotation that tells us how the original variables (foods) contributes to the PC’s (our new axis)

```
ggplot(pca$rotation) +aes(PC1, rownames(pca$rotation))+ geom_col()
```

